

SUMMARY

Applied-AI deployment lead with eight years in enterprise customer success at GLG, running a \$14M+ B2B portfolio at 100% retention as the AI person on a non-AI team. Takes AI from discovery to adoption: scopes the workflow, builds the first version on the Anthropic API, proves it with golden-set, ablation, shadow and injection evals, and drives the rollout. Five shipped builds with dated evals. Targeting AI deployment, solutions engineering and deployment strategist roles, remote-first.

SKILLS

DEPLOYMENT	Discovery & Scoping · Success Criteria & Value Realization · Implementation & Rollout · Enablement & Change Management · Human-in-the-Loop Design · AI Governance · QBR / MBR · Multi-Stakeholder Orchestration
BUILD	Python (pandas, scikit-learn) · Anthropic API & SDK · Structured Outputs · RAG (ChromaDB) · MCP Servers · Prompt Caching · Model Routing · Cloudflare Workers · JavaScript · Salesforce REST (OAuth) · SQL / BigQuery · dbt Cloud
EVALUATION	Golden Sets · Paired Ablations · Injection Fixtures · Shadow Runs · Code-Enforced Validators · Cost & Latency Tracking
TOOLS	Claude · Claude Code · ChatGPT · Gemini · Salesforce · Snowflake · Databricks · Tableau · GitHub

PROFESSIONAL EXPERIENCE

GLG

AUSTIN, TX · REMOTE · JUL 2018 - PRESENT

PE-BACKED GLOBAL EXPERT NETWORK · ENTERPRISE CUSTOMER SUCCESS, IMPLEMENTATION & ACCOUNT MANAGEMENT FOR PROFESSIONAL SERVICES FIRMS · B2B

Senior Manager, Service

OCT 2023 - PRESENT · SENIOR TEAM LEADER UNTIL THE ORG FLATTENED

- Own the post-sales lifecycle for a \$14M+ enterprise portfolio at 100% retention, including Deloitte, Accenture, Gartner, KPMG, AlixPartners, Kearney and Roland Berger, from scoping and success criteria through rollout, enablement, health monitoring and expansion; 5% YoY portfolio growth.
- Designed and deployed the LLM workflows a team of 5 client-facing managers runs on, across Claude, ChatGPT, Gemini and in-house tools, for account analysis, QBR prep, escalation drafting and data synthesis; cut MBR prep from 3 hours to 1 per account per month across four accounts.
- Built a Claude-ready knowledge base over a ~110-document library, with rules to cite each document and its date, quote compliance guidance verbatim, surface conflicts instead of resolving them, and report gaps instead of inventing process; in a paired test on 10 simple requests it answered all 10 accurately and cut an average of three back-and-forth turns.
- Running the pilot with a 60-person team through September 2026, with rollout to four more pods (about 240 people) in October gated on no increase in back-and-forth against the paired baseline.
- Contributed to AI enablement as GLG scaled in-house tooling; the main tool raised deliverables produced within the first hour from 10% to ~25% in one quarter.
- Built the team SharePoint, used daily by 60+ people and the model for three other BUs, with guardrails for answer accuracy.

Team Leader

MAY 2022 - OCT 2023

- Managed 5 client solutions specialists while carrying an enterprise book, and founded an internal Center of Excellence: lifecycle playbooks, onboarding templates and Agile delivery cadences adopted org-wide, cutting new-hire ramp to revenue from 3 months to 1.
- Built renewal forecasting and pipeline tracking in Salesforce that recognized \$300K per quarter in previously untracked revenue opportunities.

Client Solutions Associate to Senior Client Solutions Manager

JUL 2018 - MAY 2022

- Promoted three times in under four years, managing complex enterprise accounts across consulting, financial services and technology.
- Built the Salesforce account-health reporting that became the team standard.

Field Discovery Broker: Requirements Triage for Field Sales

Anthropic Messages API · Salesforce REST (OAuth) · Closed-Enum JSON Contract · Human-Approved Writes · Shadow Mode · Cloudflare Worker · samievargas.com/field-discovery

- Built for a final-round case and presented to a leadership panel: turns a field sales rep's voice note into a Salesforce record with open requirements flagged before handoff, with a closed-enum contract enforced by a five-rule validator, reject-and-retry, offline capture and an idempotent upsert; a person approves every write.
- Evaluated 22 Sep 2026: 8 of 8 golden transcripts pass with 0 hard fails; 100% escalation recall and precision on 12 proposal fixtures with 0 unneeded writes; 5 of 5 slipped go-lives flagged in shadow on fixture deals; 40 of 50 injection runs unchanged; \$0.0313 per extraction; 0 invented IDs against 99 without the closed enum (29 Aug).

Guideline Assist: Live Agent Assist and Post-Call QA

Anthropic API · Prompt Caching · Validator-Enforced Next Steps · Post-Call QA · Context Ablation · Shadow Run · samievargas.com/assist

- Next-step suggestions for a live support agent from a 55-subflow guideline library in a cached prompt, a validator that rejects any step the section does not list, and a post-call QA pass.
- Evaluated on 100 frozen ABCD test chats, 24 to 26 Sep 2026: right next action on 258 of 349 (73.9%) on Sonnet 5 at \$121.83 per 1,000 chats, level with a 73.4% no-model baseline, against 50.1% at \$47.89 on Haiku 4.5, and 82.2% at \$47.23 on Gemini 3.8 Flash with an explicit context cache, held back by a 22.4 s p95; 79.5% agreement with agents' next step on a 50-chat shadow subset; the QA still flags 20 of 100 clean chats, and 2 of 10 injection fixtures moved a suggestion.

Signal: Account Intelligence for Customer Success

Anthropic API · Model Routing · Multimodal Input · Source-Weighted Prompts · Structured Outputs · Paired Ablation · samievargas.com/signal

- Reads messy account files, images and PDFs included, and writes revenue-risk and pre-call briefs: Haiku 4.5 summarizes each document and Sonnet 4.6 synthesizes, cutting a long transcript's main call from 22,976 to at most 1,012 input tokens. Cut account prep from about an hour to one minute; validated with early users, including a senior CS leader.
- Evaluated on 13 synthetic accounts, 20 runs per arm, Sep 2026: 90% risk-type recall in both arms; 0 parse failures in 520 native-contract runs; source weighting worth 24 points of economic-buyer accuracy (93% against 69%); \$0.0308 per brief (analysis call).

Brain Dump: LLM Planner for Executive-Function Support

Anthropic API · Cloudflare Workers · Strict JSON Contract · Rule-Graded Evals · Vanilla JS · samievargas.com/brain-dump

- Turns an unstructured brain dump into a plan sized to someone's energy: three levels cap what to do now and an anxious switch softens the language, a strict JSON contract sorts every item into now, later or let go, and the page re-checks it. Early users' feedback drove a single-step capture flow.
- Evaluated on a 120-plan grid, 24 Sep 2026: prompt tuning at medium effort, which the live Worker runs, cut the median sort from 14.0 s to 7.3 s and cost from \$0.0134 to \$0.0071 a plan; low effort was faster still but more than doubled the one miss the page cannot fix (7 to 16 of 120).

Life in Pixels: RAG and a Local MCP Server

Python · ChromaDB · Local Embeddings (MiniLM) · Router + Code-Enforced Citations · Golden-Set Evals · MCP Server · github.com/SamieVargas/pixels-rag

- RAG over six months of my own daily data: a router sends each question to search, filter, sum or refusal, a validator checks every cited date and number, and the same pipeline runs as a local read-only MCP server for Claude Desktop and Claude Code.
- Evaluated on 26 questions over a synthetic export, 22 Sep 2026, Haiku 4.5: 100% route accuracy, 26 of 26 valid citations, 5 of 5 unanswerable questions refused, 85% of expected facts, \$0.0062 per question.

ANALYTICS & OTHER WORK

Instacart dbt + ML · dbt Cloud, BigQuery, scikit-learn · staging-to-mart pipeline over 3.4M orders with 35 passing tests; the cited 0.60 reorder rate hides new shoppers at 0.221 and regulars at 0.670 (random forest AUC 0.9886 against 0.8566).

ATX Foodie · Python, Socrata API, folium · 21,160 City of Austin inspections; follow-ups average 84.4 against 90.9 for routine visits, and scores drift 2.1 points toward more violations by the fifteenth inspection.

CERTIFICATIONS

Anthropic AI Fluency: Claude 101 · Claude Code 101 · Agent Skills · Claude Cowork · AI Capabilities & Limitations · AI Fluency Framework

Databricks: AI Agent Fundamentals · Generative AI Fundamentals · Databricks Fundamentals · dbt Fundamentals · Snowflake Hands-On Essentials

Google: AI Professional · Advanced Data Analytics · Business Intelligence

EDUCATION

B.S.A. Biochemistry · The University of Texas at Austin